

Gairo Peralta

gairo@berkeley.edu

gairoperalta.com

(323) 434-5684

[LinkedIn.com/in/gairoperalta](https://www.linkedin.com/in/gairoperalta)

AI/ML Engineer specializing in **Large Language Models**, **production ML systems**, and **multi-modal AI applications**

EDUCATION

TECNOLOGICO DE MONTERREY <i>MA in Applied Artificial Intelligence</i>	Mexico
THE UNIVERSITY OF TEXAS AT SAN ANTONIO <i>Visiting Student Researcher (Medical AI, Volunteer)</i>	Texas
UNIVERSITY OF CALIFORNIA, BERKELEY <i>BA in Data Science</i>	California
	2022

PROFESSIONAL EXPERIENCE

BLACKPRINT TECHNOLOGIES <i>AI/ML Engineer (Full-time)</i>	Remote	2025–Present
- Built LLM-powered GeoAI platform for 2.5M+ urban blocks with React/Mapbox/Turf.js and real-time inference - Built Node.js + Google Cloud web platform for license regularization using external POI data: density gradients by cadastral block, alcohol business flags, dashboards, and PDF export		
<i>Data Scientist</i>	Berkeley, CA	2022–2023
Delivered property analytics with Google Maps geocoding (92% accuracy) and built pandas/SQL ETL pipelines ensuring 99.9% deliverable accuracy		
LAW OFFICES OF MANUEL SOLIS <i>Lead AI/ML Engineer (Consulting)</i>	Remote	2025–Present
- Led development of custom LLM agents for T-Visa declarations and cover letters, cutting turnaround from 2 hours to 15 minutes per document (8x faster), boosting productivity; automated email-to-case DB ingestion and assignment - Built meeting intelligence platform (Zoom/Teams/Meet) that turns recordings into searchable transcripts and action-item summaries, with group-based access control across org/team		

STRTGY <i>AI/ML Engineer (Full-time)</i>	Remote	2023–2025
- Architected multi-modal GeoAI system (spatial embeddings + FAISS), driving 82% improvement in location-based decision accuracy - Built production ML pipelines: Transformer+LoRA forecasting (42% accuracy gain, 60% lower latency) and enterprise RAG (Claude 3.5 Sonnet, hybrid retrieval) with 90% relevance		

PRODENSA <i>AI Engineer (Contract)</i>	Remote	2024–2025
Built OCR + fine-tuned Llama 3.3 70B document extraction pipeline (95% accuracy); deployed via FastAPI with React/Vite front-end for real-time PDF processing		

SELECTED AI/ML PROJECTS (DETAILS)

MEDICAL QA MULTI-AGENT SYSTEM (PrimeKG + Neo4j) UT SAN ANTONIO (VOLUNTEER)	2025–Present
- Developed multi-agent LLM architecture with knowledge graph integration (PrimeKG: 4M+ nodes, 33M+ edges), achieving 54.69% accuracy on medical QA (+24.24pp) - Engineered four specialized agents (II-Medical-8B, Kimi-K2-Instruct) with Neo4j for multi-hop reasoning and systematic evaluation	

MEETING INTELLIGENCE PLATFORM LAW OFFICES OF MANUEL SOLIS (CONSULTING)	2025–Present
- Captured Zoom/Teams/Meet meetings; Speechmatics transcription + GPT-5 structured summaries (bullets, deadlines) - Indexed transcripts in Pinecone for semantic search/Q&A and retrieval across meetings - Enforced group-based RBAC for access control (admin/supervisor/user) across teams	

TECHNICAL SKILLS

GenAI/LLM: RAG, agents, fine-tuning (LoRA), embeddings, LangChain, FAISS, Pinecone, Speechmatics, RBAC
LLM Platforms: OpenAI (GPT-4o/GPT-5), Claude, Gemini, Llama
ML/DL: PyTorch, TensorFlow, scikit-learn, XGBoost, LSTM, forecasting
Backend/Cloud: FastAPI, Node.js, Google Cloud, REST APIs, MLflow, Docker
Data/Geo: Python, SQL, TypeScript, pandas, NumPy, GeoPandas, Neo4j, MongoDB, Mapbox GL, Turf.js, ArcGIS